

Exercises: Principal Component Analysis (PCA)

2023-09-26

Welcome to the exercise session on Principle Component Analysis. In this class we will go through the steps of conducting PCA on a dataset relating to kidney failure.

We will be using a couple of additional packages that are required to conduct PCA in R. The first is called 'corr' and is used for running a PCA, and the second is 'corrplot', which plots the output of corr.

First, read in the data and have a look at the columns that are available.

```
kidney_data <- read.csv("../data/chronic_kidney_disease_full.csv")
```

Now we need to check whether there are NAs in our data - remember that this can break our analysis if we don't check for it first!

```
colSums(is.na(kidney_data))
```

```
##      X    age    bp    sg    al    su    rbc    pc    pcc    ba
##      0      0      0      0      0      0      0      0      0      0
##    bgr    bu    sc    sod    pot    hemo    pcv    wbcc    rbcc    htn
##      0      0      0      0      0      0      0      0      0      0
##    dm    cad  appet    pe    ane    class no_name
##      0      0      0      0      0      0      0      0
```

```
# is.na checks for NA values in each cell of the
# dataframe. colSums returns the sum of the columns - if any one entry is NA,
# then the colSums will give us an NA! So this is a neat way of checking for the
# presence of NA values.
```

Check which of the variables are categorical and which are continuous - remember that we need continuous data for PCA! Extract the columns that correspond to continuous data. How many variables do you get?

```
head(kidney_data)
```

```
##   X age bp    sg al su    rbc    pc    pcc    ba bgr bu  sc sod pot
## 1 0  48 80 1.020 1  0      ?   normal notpresent notpresent 121 36 1.2 ?  ?
## 2 1   7 50 1.020 4  0      ?   normal notpresent notpresent  ? 18 0.8 ?  ?
## 3 2  62 80 1.010 2  3 normal   normal notpresent notpresent 423 53 1.8 ?  ?
## 4 3  48 70 1.005 4  0 normal abnormal   present notpresent 117 56 3.8 111 2.5
## 5 4  51 80 1.010 2  0 normal   normal notpresent notpresent 106 26 1.4 ?  ?
## 6 5  60 90 1.015 3  0      ?      ? notpresent notpresent  74 25 1.1 142 3.2
##   hemo pcv wbcc rbcc htn  dm cad appet  pe ane class no_name
## 1 15.4  44 7800  5.2 yes yes  no  good  no  no   ckd
## 2 11.3  38 6000   ?  no  no  no  good  no  no   ckd
## 3  9.6  31 7500   ?  no yes  no  poor  no yes   ckd
## 4 11.2  32 6700  3.9 yes  no  no  poor yes yes   ckd
## 5 11.6  35 7300  4.6  no  no  no  good  no  no   ckd
## 6 12.2  39 7800  4.4 yes yes  no  good yes  no   ckd
```


```
# Select the continuous variables
data_for_PCA <- kidney_data %>% select(age:sg, bgr:rbcc)
```



```

# Correlation matrix plot
corrplot(
  corr_matrix,
  # Only plot the lower triangular
  type = 'lower',
  # Visualisation method: show ellipses
  # eccentricity scaled to the correlation value
  method = 'ellipse',
  # Sort by angular order of the eigenvectors
  order = 'AOE',
  # Rotate the text labels
  tl.srt = 45
)

```



Now, we can perform the PCA. We can use the command `prcomp` to perform the PCA on the scaled data.

```

# Perform PCA
PCA <- prcomp(corr_matrix)

# SD of the principal components
sdevs <- PCA$sdev

# Total variance
total_var <- sum(sdevs^2)

# Get proportion variance (divide SD^2 by total variance)
proportion_of_variance <- data.frame(Proportion_of_Variance = (sdevs^2) / total_var * 100)

```

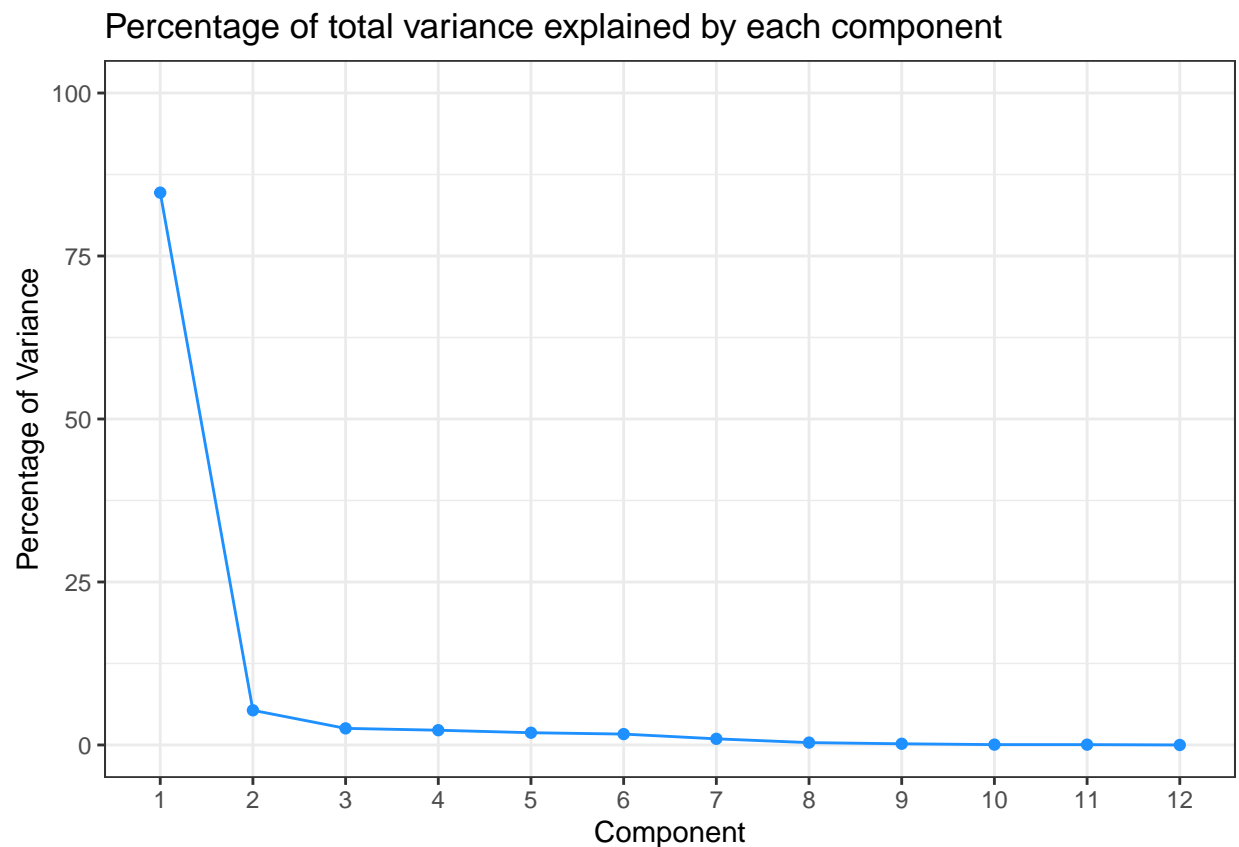
```

proportion_of_variance$Component <- as.character(row.names(proportion_of_variance))

# Scree plot
scree_plot <- ggplot(
  proportion_of_variance, aes(
    fct_reorder(Component, -Proportion_of_Variance), Proportion_of_Variance, group = 1)
) +
  # Add line plot
  geom_line(color = 'dodgerblue') +
  # Add scatter plot
  geom_point(color = 'dodgerblue') +
  # Change y limit, xy labels, title, and theme
  ylim(0, 100) +
  xlab('Component') +
  ylab('Percentage of Variance') +
  ggtitle('Percentage of total variance explained by each component') +
  theme_bw()


scree_plot

```




We can also use `fviz_pca_var` to plot the each variable in terms of the principal components.

```
fviz_pca_var(PCA, col.var = 'black')
```



Finally, we can use a cos2 plot to transform the principal components back into our original:

```
fviz_cos2(PCA, choice = "var", axes = 1)
```



We can also use PCA to predict whether or not a patient will have chronic kidney disease. We split the data into a train set, on which we perform the PCA, and a test set, for which we try to predict whether each patient has chronic kidney disease.

We can then plot the principal components for each of these patients on a scatter plot, and see whether the patients are separated by their kidney condition.

```
# Take the continuous data
data_for_prediction <- kidney_data %>% select(X, age:sg, bgr:rbcc)

# Remove NAs
data_for_PCA <- data_for_prediction %>%
  mutate_if(is.character, as.numeric) %>% na.omit()
```

```
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
```

```
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
```

```
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
```

```
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
```

```
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
```

```
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
```

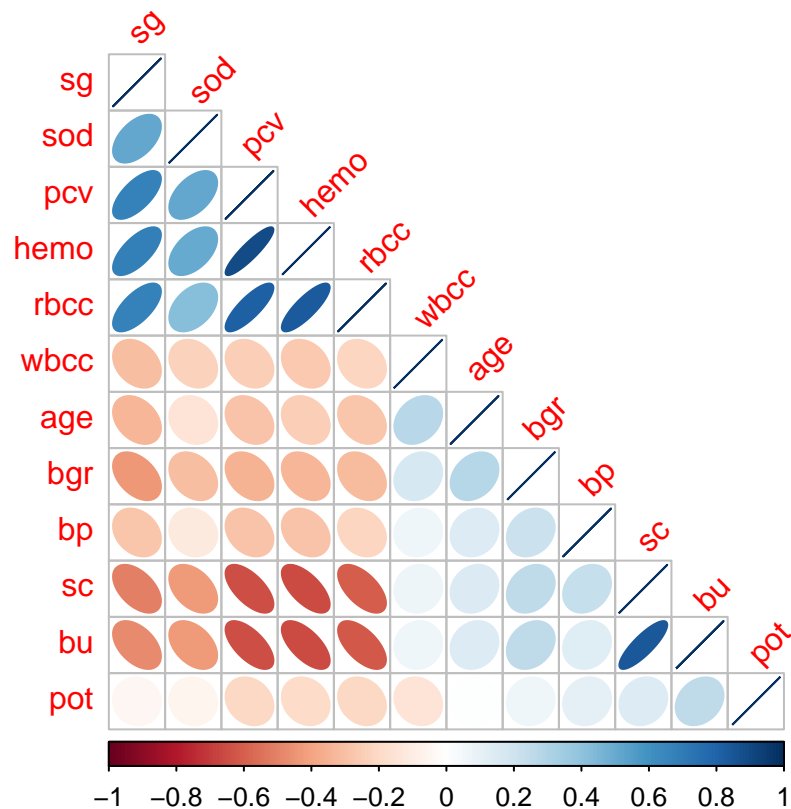
```
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
```

```
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion
## Warning in mask$eval_all_mutate(quo): NAs introduced by coercion

# Add the category labels, i.e. does each person have chronic kidney disease
data_for_PCA$class <- kidney_data %>% filter(X %in% data_for_PCA$X) %>%
  select(class)

# Split the data into testing and training - choose a random sample from the full
# dataset
train <- data_for_PCA %>% sample_frac(0.70)
test  <- dplyr::anti_join(data_for_PCA, train, by = 'X')

#Scale the training set and produce a correlation matrix
scale_train <- as.data.frame(scale(train %>% select(age:sg, bgr:rbcc)))
corr_matrix <- cor(scale_train)
corrplot(corr_matrix, type = 'lower', method = 'ellipse', order = 'AOE',
          tl.srt = 45)
```



```

# Perform the PCA on the scaled training set
PCA <- prcomp(scale_train, scale. = TRUE, center = TRUE)

# Scale and perform PCA on the test set
scale_test <- as.data.frame(scale(test %>% select(age:sg, bgr:rbcc)))
pred <- predict(PCA, newdata = scale_test)
train_df <- as.data.frame(PCA$x[, 1:2])
train_df$Diagnosis <- unname(unlist(train$class))

# Transform the testing data using the principal components.
predict_df <- as.data.frame(pred[, 1:2])
predict_df$Diagnosis <- unname(unlist(test$class))

# Plot the results
scatter <- ggplot(predict_df, aes(x = PC1, y = PC2, color = Diagnosis)) +
  geom_point() + theme_bw()
scatter

```

